

Software Engineering



Submitted by:

Muhammad Awais (FA23-BCS-103)

Section: FA23-BCS_B

Submitted on:

October 11, 2025

Department of Computer Science

COMSATS University Islamabad

Database Schema (MongoDB)

We are using MongoDB because the data is mostly read-heavy, large, and semi-structured.

Main Collections:

books

Fields:

- _id (ObjectId)
- book_id (Integer)
- goodreads_book_id (Integer)
- title (String)
- authors (String)
- original_publication_year (Integer)
- average_rating (Float)
- ratings_count (Integer)
- image_url (String)
- small_image_url (String).

Indexes Used:

Indexes improve reading speed.

Books Collection Indexes:

1. **{title: 1, authors: 1}**
Used for searching by title and author.
2. **{average_rating: -1}**
Used to quickly get books with highest ratings.
3. **{book_id: 1}**
Used for direct book lookups.

Ratings Collection Indexes:

1. **{book_id: 1}**
Used for finding all ratings of a book.
2. **{user_id: 1, book_id: 1}** (unique)
Used for user-specific book ratings.

Tags Collection Indexes:

1. **{tag_id: 1}**
Used for tag lookups.
2. **{tag_name: 1}**
Used for tag name searches.

Book Tags Collection Indexes:

1. **{tag_id: 1}**
Used for finding books by tag.
2. **{goodreads_book_id: 1}**
Used for finding tags by book.

To-Read Collection Indexes:

1. **{user_id: 1, book_id: 1}** (unique)
Used for user's reading list.

These indexes reduce query time but increase storage a little.

Why these Indexes?

- Title and author search becomes fast
- Rating queries return faster
- High-rating sorting becomes efficient
- User-specific operations are optimized
- Better response performance for REST API
- Tag-based filtering works efficiently

Without indexes, MongoDB would scan the whole database every time.

Trade-offs / Limitations:

Pros:

- Fast read performance for common queries
- Easy scaling with sharding
- Flexible schema for evolving data
- Good for complex aggregations
- Efficient for high-volume reads

Cons:

- Indexes take extra storage
- Write performance becomes slightly slower
- Data is normalized across collections
- Some queries require joins between collections
- More complex query patterns needed for related data